

finalquestion

```
library(itsmr)
library(fpp2)


## Loading required package: ggplot2
## Loading required package: forecast
##
## Attaching package: 'forecast'
## The following object is masked from 'package:itsmr':
##
##   forecast
## Loading required package: fma
##
## Attaching package: 'fma'
## The following objects are masked from 'package:itsmr':
##
##   airpass, strikes
## Loading required package: expsmooth
library(expsmooth)
library(forecast)
library(tibbletime)

##
## Attaching package: 'tibbletime'
## The following object is masked from 'package:stats':
##
##   filter
library(tidyverse)


## -- Attaching packages -----
## v tibble  2.1.3      v dplyr   0.8.4
## v tidyr   1.0.2      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0
## v purrr   0.3.3
## -- Conflicts -----
## x dplyr::filter() masks tibbletime::filter(), stats::filter()
## x dplyr::lag()    masks stats::lag()
library(tsbox)
library(knitr)

##(a)
beer_dat=dget("beer.Rput")
# remove the last 12 values
beer<-head(beer_dat,-12)
```


```
# plot the data  
autoplot(beer)+ggtitle("Beer")
```



```
ggAcf(beer)
```



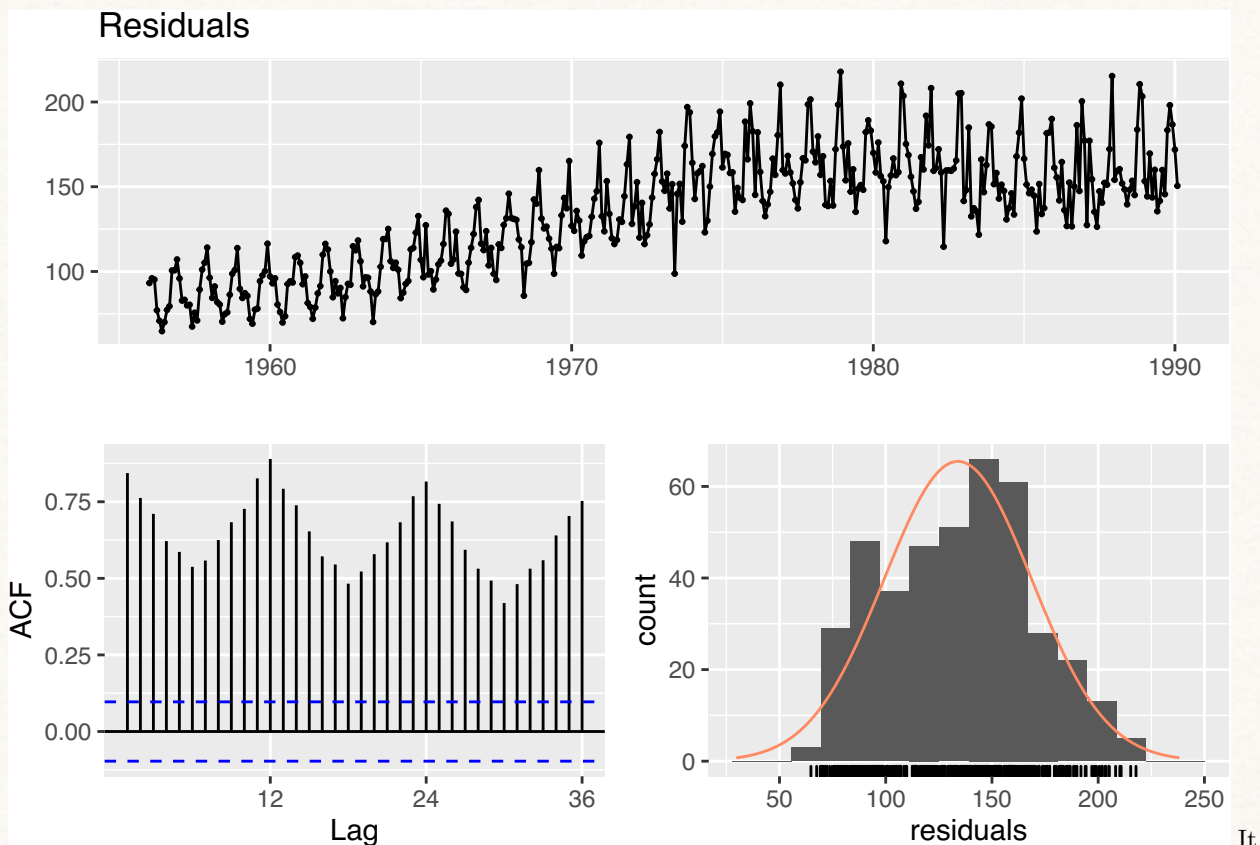
```
ggPacf(beer)+ggtitle("Beer")
```

Lag It shows that the data is not stationary but contain the seasonal component and noise. Also the above ACF plot shows the seasonal component and the data are correlated. It is obvious that this dataset is not a white noise process.

```
# check the residual by deseasonalizing
#sales_ma5<-ma(beer,order=5)
#tslm<-tslm(beer~trend+I(trend^2)+season)
#stl<-stl(beer,s.window=12)
#checkresiduals(stl$time.series[, 'remainder'])
checkresiduals(beer)
```

```
## Warning in modeldf.default(object): Could not find appropriate degrees of
## freedom for this model.
```



It shows that the data is not stationary but contain the seasonal component and noise. Also the above ACF plot and the residual plot shows the seasonal component and the data are correlated. It is obvious that this dataset is not a white noise process.

find the arima model

```
beer_ari<-auto.arima(beer,stepwise=FALSE, seasonal=FALSE, ic='aic',trace=TRUE)
```

```
##
## Fitting models using approximations to speed things up...
##
## ARIMA(0,1,0) : 3579.226
## ARIMA(0,1,0) with drift : 3581.204
## ARIMA(0,1,1) : 3547.203
## ARIMA(0,1,1) with drift : 3549.134
## ARIMA(0,1,2) : 3493.171
## ARIMA(0,1,2) with drift : 3489.323
## ARIMA(0,1,3) : 3490.168
## ARIMA(0,1,3) with drift : 3486.407
## ARIMA(0,1,4) : 3468.429
## ARIMA(0,1,4) with drift : 3464.997
## ARIMA(0,1,5) : 3469.946
## ARIMA(0,1,5) with drift : 3466.611
## ARIMA(1,1,0) : 3556.574
## ARIMA(1,1,0) with drift : 3558.536
## ARIMA(1,1,1) : 3483.419
## ARIMA(1,1,1) with drift : 3480.55
## ARIMA(1,1,2) : 3485.297
## ARIMA(1,1,2) with drift : 3482.359
```

```

## ARIMA(1,1,3) : 3486.94
## ARIMA(1,1,3) with drift : 3483.973
## ARIMA(1,1,4) : 3471.249
## ARIMA(1,1,4) with drift : 3468.278
## ARIMA(2,1,0) : 3547.96
## ARIMA(2,1,0) with drift : 3549.898
## ARIMA(2,1,1) : 3540.196
## ARIMA(2,1,1) with drift : 3542.157
## ARIMA(2,1,2) : 3481.216
## ARIMA(2,1,2) with drift : 3477.923
## ARIMA(2,1,3) : 3482.351
## ARIMA(2,1,3) with drift : 3479.135
## ARIMA(3,1,0) : 3549.09
## ARIMA(3,1,0) with drift : 3551.011
## ARIMA(3,1,1) : 3543.732
## ARIMA(3,1,1) with drift : 3545.695
## ARIMA(3,1,2) : 3481.842
## ARIMA(3,1,2) with drift : 3476.954
## ARIMA(4,1,0) : 3541.237
## ARIMA(4,1,0) with drift : 3543.099
## ARIMA(4,1,1) : 3449.546
## ARIMA(4,1,1) with drift : 3443.193
## ARIMA(5,1,0) : 3543.721
## ARIMA(5,1,0) with drift : 3545.553
##
## Now re-fitting the best model(s) without approximations...
##
##
##
## Best model: ARIMA(4,1,1) with drift

```

Now analyze the model:

```
summary(beer_ari)
```

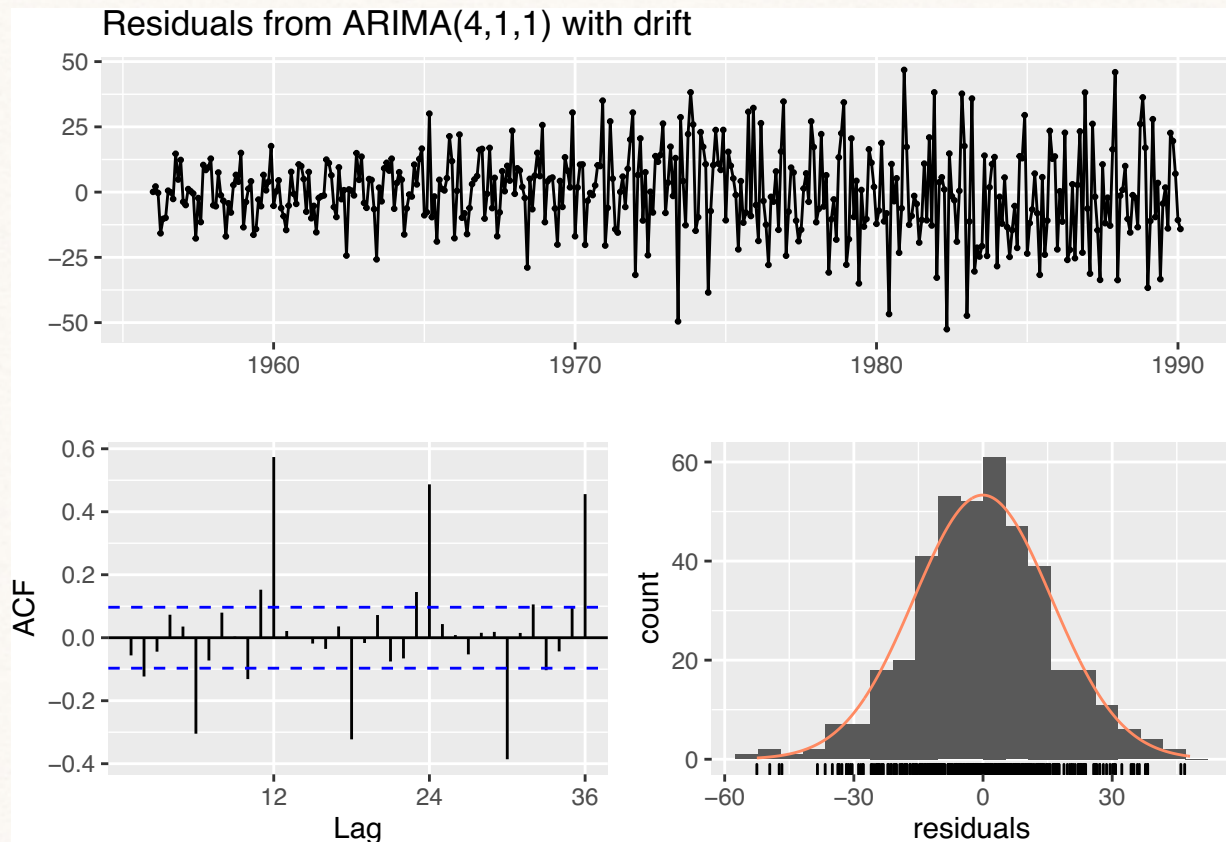
```

## Series: beer
## ARIMA(4,1,1) with drift
##
## Coefficients:
##          ar1      ar2      ar3      ar4      ma1      drift
##          0.4462  0.0028  0.0763 -0.3170 -0.9395  0.1977
## s.e.  0.0477  0.0518  0.0517  0.0475  0.0127  0.0632
##
## sigma^2 estimated as 261.2:  log likelihood=-1716.54
## AIC=3447.08  AICc=3447.36  BIC=3475.18
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.0588431 16.02236 12.37635 -1.226062 9.171119 1.306263
##              ACF1
## Training set -0.05614064

```



```
checkresiduals(beer_ari)
```



```
##
## Ljung-Box test
##
## data: Residuals from ARIMA(4,1,1) with drift
## Q* = 377.63, df = 18, p-value < 2.2e-16
##
## Model df: 6. Total lags used: 24
#test(checkresiduals(beer_ari))
```

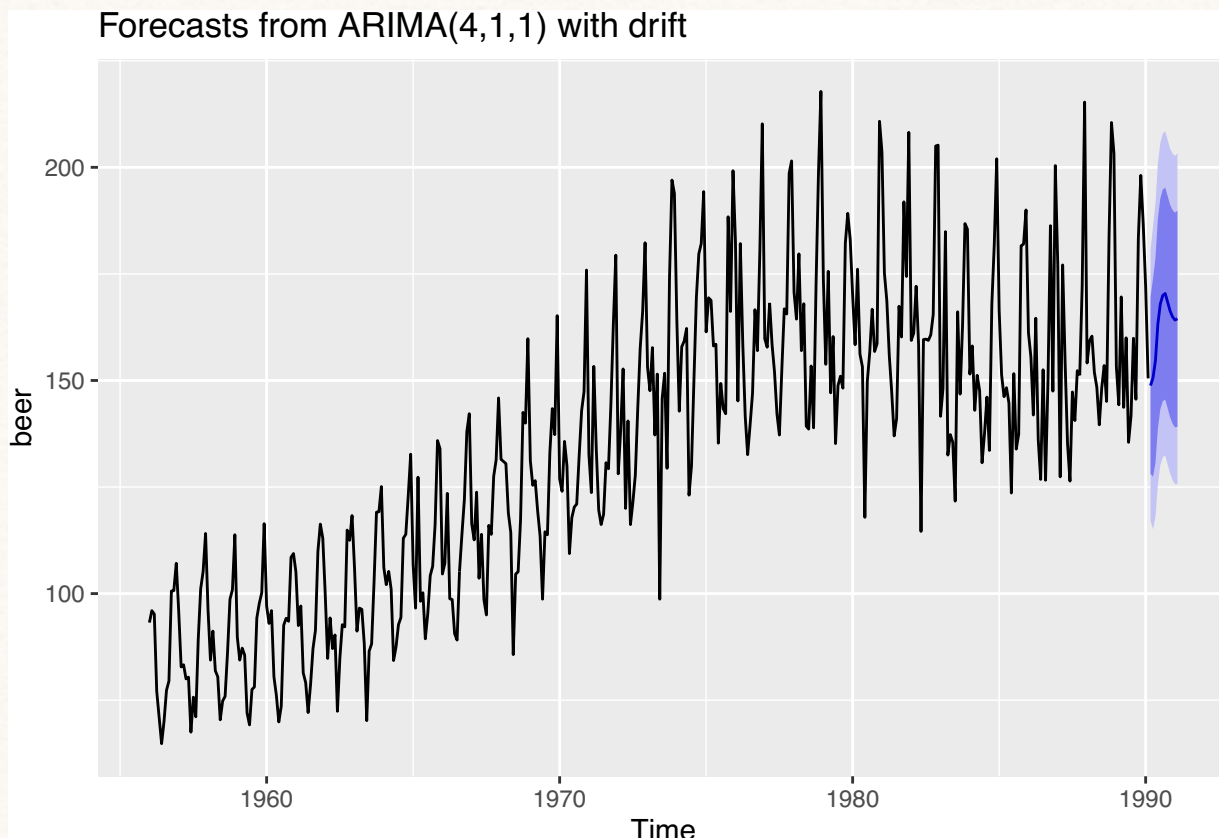
Construct the 95% confidence interval. Showing the bounds for component θ and ϕ

```
confint(beer_ari)
```

```
##           2.5 %      97.5 %
## ar1    0.35276931  0.5396158
## ar2   -0.09861421  0.1043077
## ar3   -0.02492603  0.1775627
## ar4   -0.41008956 -0.2239423
## ma1   -0.96429447 -0.9146189
## drift  0.07376610  0.3216664
```

Now forecast on the 12 removed data

```
beer_forecast<-forecast(beer_ari,h=12)
autoplot(beer_forecast)
```



Here's the result of prediction presented in table:

```
bt<-log(tail(beer_dat,12))
error<-bt-beer_forecast$mean
data<-data.frame(error, beer_forecast$lower, beer_forecast$upper,beer_forecast$mean)
colnames(data)<-c('error', '2.5%', '97.5%', 'predict')
data
```


##	error	2.5%	97.5%	predict	NA	NA
## 1	-143.7117	128.0945	117.1307	169.5165	180.4803	148.8055
## 2	-145.5951	127.4111	115.1201	173.8479	186.1389	150.6295
## 3	-149.6245	130.6737	117.9794	178.6337	191.3279	154.6537
## 4	-158.3575	138.6542	125.6256	187.8773	200.9059	163.2658
## 5	-162.9536	143.2583	130.1858	192.6473	205.7197	167.9528
## 6	-164.9548	145.1331	131.9937	194.7749	207.9143	169.9540
## 7	-165.5043	145.5410	132.3823	195.2557	208.4145	170.3984
## 8	-163.1195	143.4968	130.3211	193.2759	206.4516	168.3864
## 9	-160.9738	141.3979	128.2084	191.2287	204.4182	166.3133
## 10	-159.6554	139.9032	126.6503	189.9740	203.2269	164.9386
## 11	-159.0816	139.0031	125.6744	189.3598	202.6884	164.1814
## 12	-159.4786	139.1259	125.7065	189.8257	203.2451	164.4758

##(c) First deseasonalizing

```
beer_linear<- tslm(beer~trend)
beer_quadratic<- tslm(beer~trend+I(trend^2))
beer_ma3<-ma(beer,order=3)
beer_ma5<-ma(beer,order=5)
#plot the fitted graph
```



```
beer_with_fit<-cbind(beer,Linear_trend=fitted(beer_linear),Quadratic_trend=fitted(beer_quadratic),MA3=fitted(beer_ma3),MA5=fitted(beer_ma5))
autoplot(beer_with_fit)
```




The residual plots by removing these the quadratic trends respectively from the dataset :

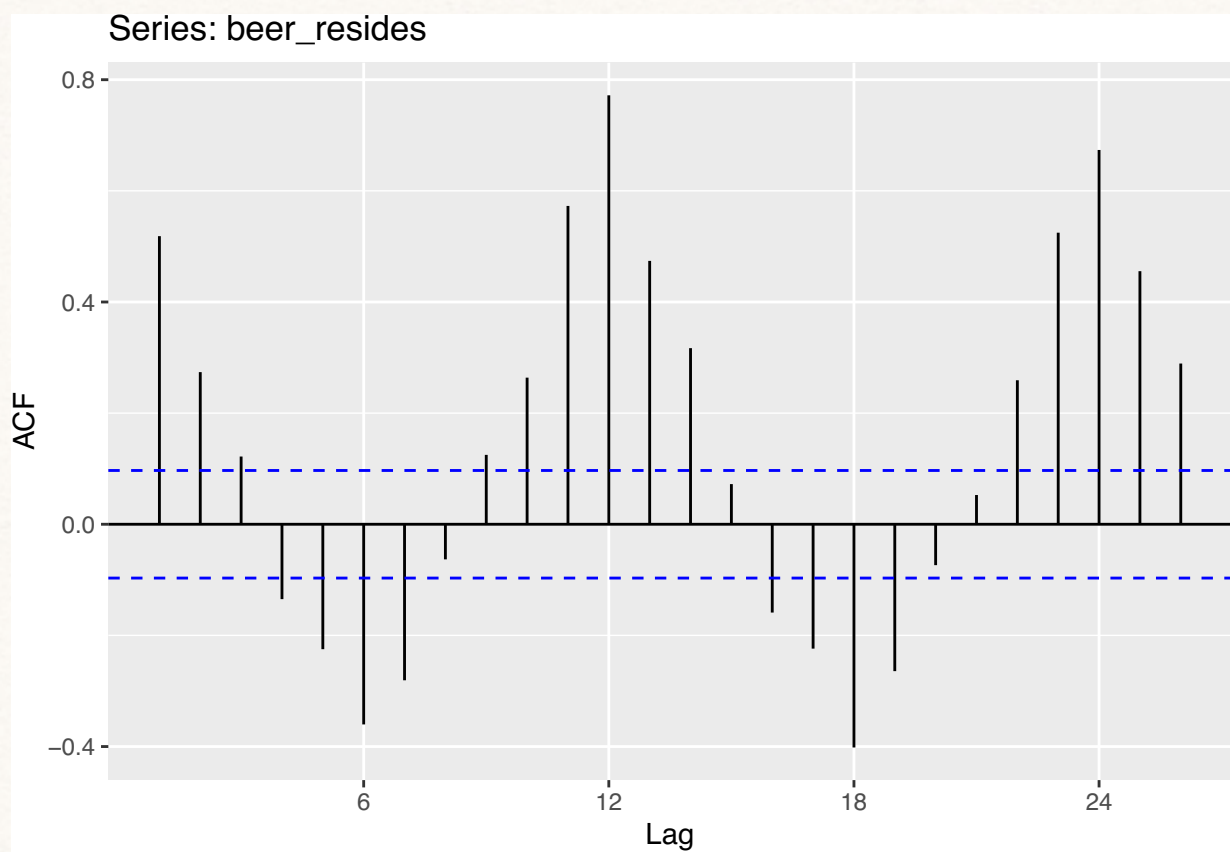
```
beer_resides<-cbind(
  res_quad=beer-fitted(beer_quadratic))
```

```
autoplot(beer_resides,facet=TRUE)
```

```
## Warning: Ignoring unknown parameters: facet
```



```
#autoplot(sales_resides, facet=TRUE)
ggAcf(beer_resides)
```




```
frequency(beer_resides)
```

```
## [1] 12
```

```
beer_seasonlm<-tslm(beer_resides~season)
```


```
autoplot(fitted(beer_seasonlm))
```




```
beers_resids<-cbind(res_ma5=beer_resides,
```

```
                  res_ma5_season=beer_resides-fitted(beer_seasonlm))
```


```
autoplot(beers_resids,facet=TRUE)
```


```
beer_deseason<-beer-mean(beer)-fitted(beer_seasonlm)
autoplot(beer_deseason)
```



```
beer_qua<- tslm(beer_deseason~trend+I(trend^2))  
beer_dd<-beer_deseason-fitted(beer_qua)  
autoplot(beer_dd)
```



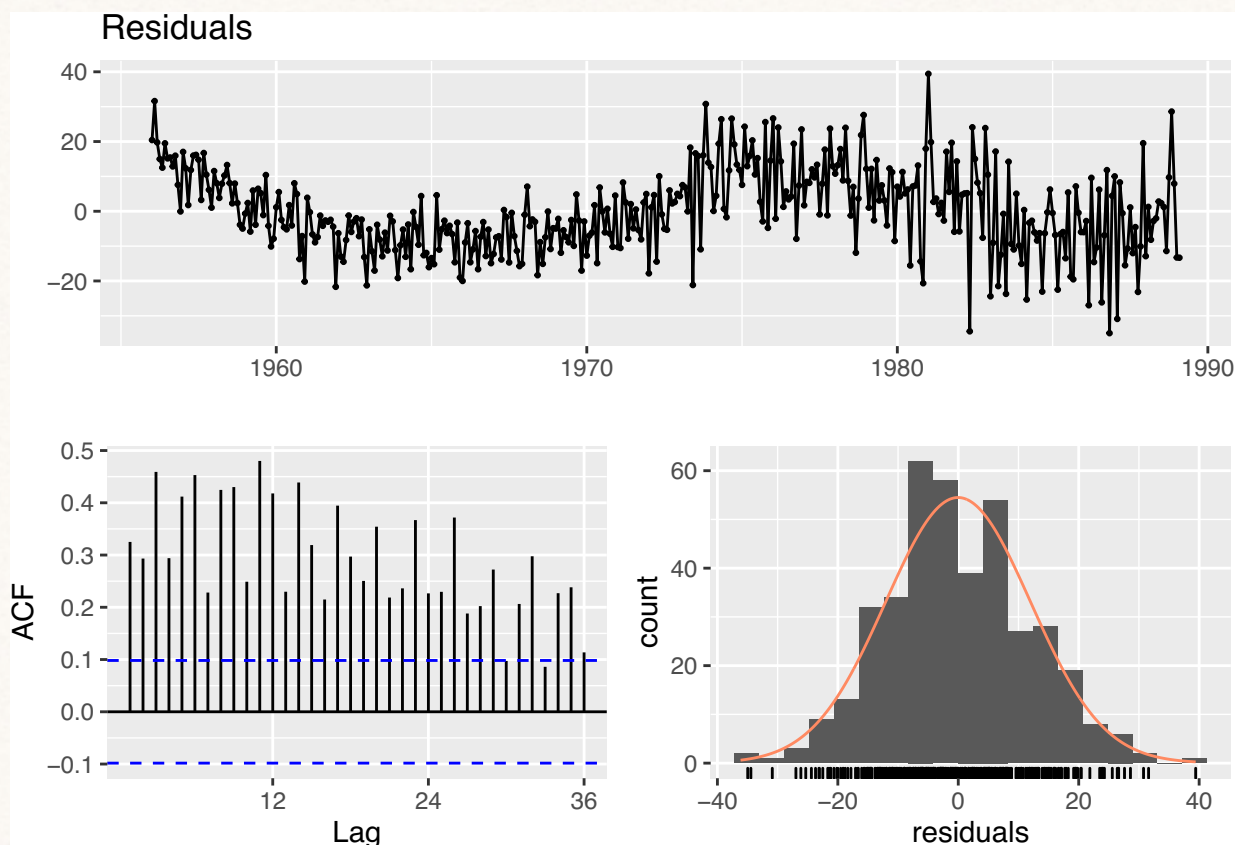
```
ggAcf(beer_dd)
```



By the autoplot and residual plot of the data, it shows that the data may possibly be a white noise process, with low correlation.

```
beer_ddd<-head(beer_dd,-12)
checkresiduals(beer_ddd)
```

```
## Warning in modeldf.default(object): Could not find appropriate degrees of
## freedom for this model.
```

Try to do the ARIMA model again:

```
pro<-auto.arima(beer_ddd,stepwise=FALSE, seasonal=FALSE, ic='aic',trace=TRUE)
```

```
##
## Fitting models using approximations to speed things up...
##
## ARIMA(0,1,0) : 3215.129
## ARIMA(0,1,0) with drift : 3217.114
## ARIMA(0,1,1) : 2946.174
## ARIMA(0,1,1) with drift : 2947.363
## ARIMA(0,1,2) : 2939.357
## ARIMA(0,1,2) with drift : 2940.609
## ARIMA(0,1,3) : 2924.046
## ARIMA(0,1,3) with drift : 2925.451
## ARIMA(0,1,4) : 2925.712
## ARIMA(0,1,4) with drift : 2927.091
## ARIMA(0,1,5) : 2915.388
## ARIMA(0,1,5) with drift : 2916.907
## ARIMA(1,1,0) : 3111.811
## ARIMA(1,1,0) with drift : 3113.748
## ARIMA(1,1,1) : 2952.66
## ARIMA(1,1,1) with drift : 2953.785
## ARIMA(1,1,2) : 2953.807
## ARIMA(1,1,2) with drift : 2954.976
## ARIMA(1,1,3) : 2933.006
## ARIMA(1,1,3) with drift : 2934.259
## ARIMA(1,1,4) : 2933.831
```

```
## ARIMA(1,1,4)          with drift          : 2935.139
## ARIMA(2,1,0)          : 3005.483
## ARIMA(2,1,0)          with drift          : 3007.368
## ARIMA(2,1,1)          : 2929.203
## ARIMA(2,1,1)          with drift          : 2930.578
## ARIMA(2,1,2)          : 2928.436
## ARIMA(2,1,2)          with drift          : 2929.725
## ARIMA(2,1,3)          : 2910.104
## ARIMA(2,1,3)          with drift          : 2912.099
## ARIMA(3,1,0)          : 2998.079
## ARIMA(3,1,0)          with drift          : 2999.973
## ARIMA(3,1,1)          : 2928.147
## ARIMA(3,1,1)          with drift          : 2929.609
## ARIMA(3,1,2)          : 2927.022
## ARIMA(3,1,2)          with drift          : 2928.434
## ARIMA(4,1,0)          : 2959.86
## ARIMA(4,1,0)          with drift          : 2961.755
## ARIMA(4,1,1)          : 2915.703
## ARIMA(4,1,1)          with drift          : 2917.263
## ARIMA(5,1,0)          : 2934.629
## ARIMA(5,1,0)          with drift          : 2936.461
```

```
## Now re-fitting the best model(s) without approximations...
```

```
##
```

```
##
```

```
##
```

```
##
```

```
## Best model: ARIMA(2,1,3)
```

```
summary(pro)
```

```
## Series: beer_ddd
```

```
## ARIMA(2,1,3)
```

```
##
```

```
## Coefficients:
```

```
##          ar1      ar2      ma1      ma2      ma3
```

```
##          0.0996  0.4361 -1.1228 -0.4402  0.6487
```

```
## s.e.    0.1378  0.1212  0.1208  0.2098  0.1093
```

```
##
```

```
## sigma^2 estimated as 89.82:  log likelihood=-1455.05
```

```
## AIC=2922.11  AICc=2922.33  BIC=2946.01
```

```
##
```

```
## Training set error measures:
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
```

```
## Training set -0.202667  9.405641  7.28973  50.03155  318.183  0.7778955 -0.03251424
```

```
confint(pro)
```

```
##          2.5 %      97.5 %
```

```
## ar1 -0.1705255  0.36979818
```

```
## ar2  0.1985049  0.67371573
```

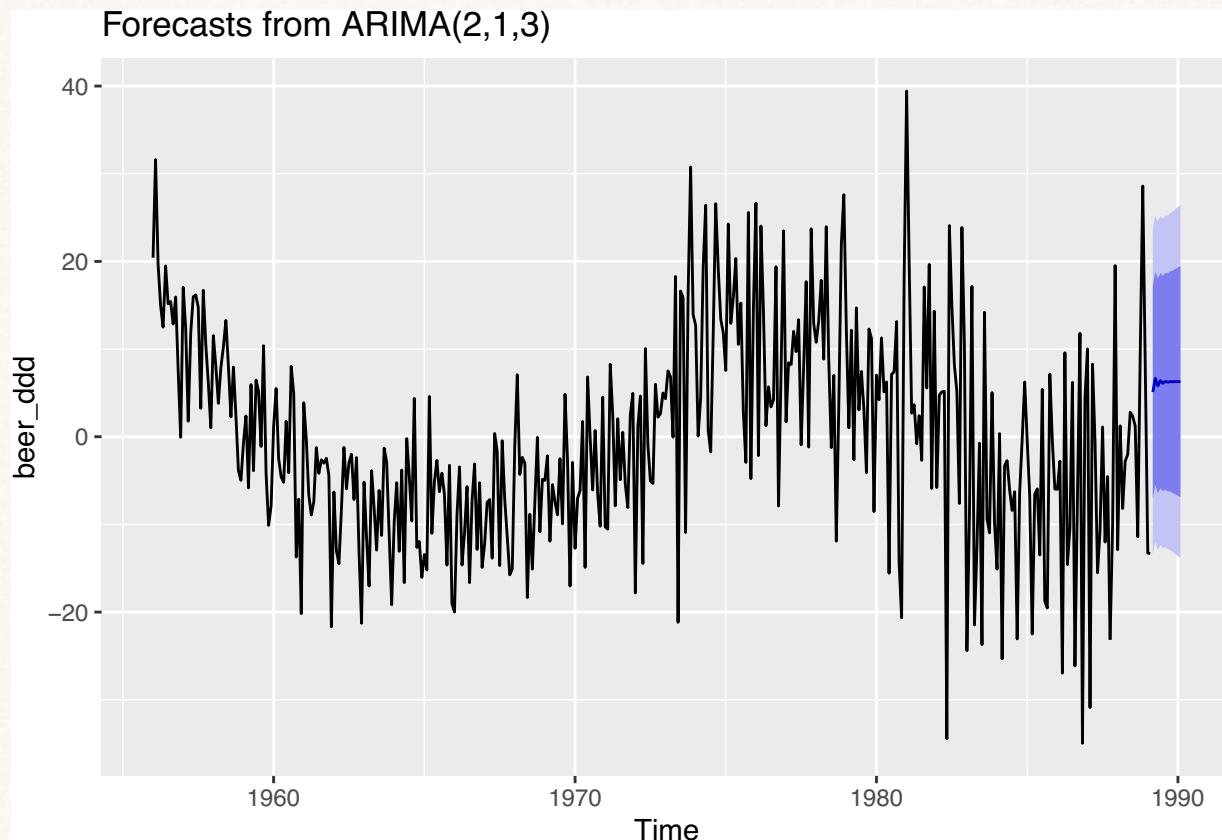
```
## ma1 -1.3595752 -0.88611656
```

```
## ma2 -0.8514258 -0.02906865
```

```
## ma3  0.4344729  0.86291610
```

Here the 95% confidence interval shows the bounds of parameter.

```
# forecast the 12-step data
foo<-forecast(pro,h=12)
autoplot(foo)
```



Now show the result by table

```
#b1<-log(tail(beer_dd,12))
b1<-tail(beer_dd,12)
e1<-b1-foo$mean
d1<- data_frame(e1,foo$lower,foo$upper,foo$mean)
```

```
## Warning: `data_frame()` is deprecated, use `tibble()`.
## This warning is displayed once per session.
```

```
colnames(d1)<-c('error','2.5%','97.5%','prediction')
d1
```

```
## # A tibble: 12 x 4
##   error `2.5%` `97.5%` prediction
##   <dbl>   <dbl>   <dbl>   <dbl>
## 1  -3.62   -7.05  -13.5    17.2
## 2 -17.2    -5.49  -11.9    18.8
## 3   4.24   -6.44  -12.9    18.1
## 4  -7.21   -5.87  -12.4    18.7
## 5  -9.76   -6.19  -12.7    18.4
## 6   1.83   -6.04  -12.6    18.7
## 7 -16.6    -6.23  -12.8    18.6
## 8   3.69   -6.27  -12.9    18.9
## 9  10.5    -6.44  -13.2    19.0
```


## 10	-14.4	-6.56	-13.4	19.1	26.0	6.29
## 11	-0.328	-6.73	-13.6	19.3	26.2	6.28
## 12	-12.8	-6.89	-13.9	19.5	26.5	6.29

It is obvious that comparing with part b, the error of prediction become much smaller.